

Reinforcement Learning – COMS4047A

Hierarchical RL


Prof. Benjamin Rosman

Benjamin.Rosman1@wits.ac.za / benjros@gmail.com


Outline so far

- We started with **dynamic programming**
 - But want to solve problems when we **don't know a model**
- So we used **Monte Carlo** methods
 - But this needs full **episode** returns
- So we moved to tabular **temporal difference** methods
 - But these don't scale to **large problems**
- So we introduced **function approximation**
 - So we're good, right?


Some really hard problems!




(a) Hopper-v1




(b) Walker2d-v1



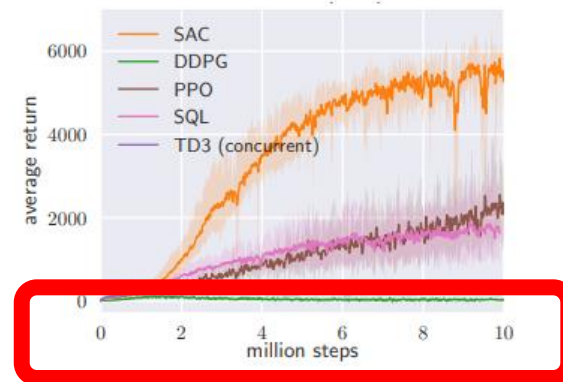
(c) HalfCheetah-v1



(d) Ant-v1



(e) Humanoid-v1





(f) Humanoid (rllab)

Soft actor-critic, Haaronoja et al 2018

Why is RL hard?

- **Sparse**
 - Most actions give **no reward feedback**
- **Delayed**
 - Rewards may come after executing whole trajectories



How does RL work?

- Solves a “flat” problem!
- Want a robot to make coffee?
- +1 for making coffee, -1 otherwise
- Find a set of actions to achieve this
 - What are those actions? Motor commands? Complex programs?

The most important task

- How do you make coffee?

What about making tea?

*Fill Kettle
Plug in Kettle*

...

Boil water

Add coffee to cup

Pour water into cup

Stir liquid

Enjoy deliciousness

Hierarchical decomposition

- Decompose the problem into **smaller** ones
- **Solve** each of those smaller problems
 - Maybe recursively decompose the subproblems?
- Exploits structure in problems
 - Maybe we can **reuse** these subproblems elsewhere?

Some questions

$$M = (S, A, T, R, \gamma)$$

- What is the **best way** of building a hierarchy?
- Can we **learn** it?
- What **kind of hierarchies** can we construct?
 - **State and action abstraction**


Hierarchical RL

- RL typically solves a *single* problem *monolithically*
- Hierarchical RL:
 - Create and use higher-level **macro-actions** (skills)
 - Problem now contains subproblems
 - Each **subproblem** is also an **RL problem**
- Hierarchies of abstract machines (HAMs), MAXQ, **Options**

Skill hierarchies

- **Hierarchical RL**: base hierarchical control on **skills**
 - Component of behaviour
 - Performs continuous, low-level control
 - Can treat as discrete action
- Behaviour is **modular** and compositional
 - Like **functions** in a program!

```
def kick_ball(self, dest):  
    # bunch of low level actions here
```



Options [Sutton, et al 2000]


- Theoretical basis for skill acquisition,
 - learning and planning using **higher-level actions** (options)
- An option is a **temporally-extended** action
 - An action executed over many timesteps

 *Like a policy*

Intuitively

- An option o is a policy unit:

- Initiation set
- Termination condition
- Option policy




Formally

- A (Markov) option o is defined by
 - **Initiation set**: $I_o \subseteq S$
 - **Policy**: $\pi_o: S \times A \rightarrow [0, 1]$
 - **Termination condition**: $\beta_o: S \rightarrow [0, 1]$
- Can have non-Markov options
 - Functions not solely of state, but also **execution history**
 - Run for at most n steps, repeat n times, etc
 - Not often used, but can be useful


What about actions?

- **Primitive** action a can be represented by an option:
- $I_a = S$
- $\pi_a(s, b) = \begin{cases} 1, & a = b \\ 0, & a \neq b \end{cases}$
- $\beta_a(s) = 1 \ \forall s \in S$
- A primitive action can be executed **anywhere**, lasts exactly **one timestep**, and always chooses **action a**

Options as actions



Option



Set of decisions available

Questions...

- Given an MDP (S, A, R, T, γ)
 - Replace A with set of options O (some may be primitive actions)
- How do we characterise the resulting problem?
- How do we plan with options?
- How do we learn with options?
- How do we characterise resulting policies?

SMDPs

- Resulting problem is **Semi-Markov** Decision Process
 - (S, O, T, R, γ)
- S is the set of **states**
- O is the set of **options**
- $T = \Pr(s', t \mid o, s)$ is the **transition** model
- $R(s, o, s', t)$ is the **reward** function
- γ is the **per-step** discount factor

SMDPs

- Note
 - All times are integers
 - “Semi” means transitions can last $t > 1$ timesteps
 - Transition and reward functions involve time taken for option to execute

Planning with options

- Regular Bellman equation:

$$Q_{\pi}(s, a) = \sum_{s'} p(s'|s, a) [r(s, a, s') + \gamma Q_{\pi}(s', \pi(s'))]$$

- Bellman equation with options:

$$\boxed{Q_{\pi}(s, o)} = \sum_{s', t} p(s', t | s, o) \left[\boxed{r(s, o, s', t)} + \gamma^t \boxed{Q_{\pi}(s', \pi(s'))} \right]$$

Value of o in s *Immediate reward* *Expected future value*


Learning and Planning

- For learning:
 - Collect stochastic samples
 - Use SMDP Bellman equation
- For planning:
 - Synchronous Value Iteration
 - Value Iteration using the SMDP Bellman Equation


$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

$$Q(S_t, O_t) = Q(S_t, O_t) + \alpha \left[\sum_{i=1}^k \gamma^{i-1} R_{t+i} + \gamma^k \max_o Q(S_{t+k}, o) - Q(S_t, O_t) \right]$$


Example



*4 stochastic
primitive actions*



*8 multi-step options
(to each room's 2 hallways)*






Target
Hallway




(Sutton, Precup and Singh, AIJ 1999)

Example

Primitive
options
 $\mathcal{O}=\mathcal{A}$



Hallway
options
 $\mathcal{O}=\mathcal{H}$



Initial Values


Iteration #1

Iteration #2


Example

Primitive
and
hallway
options


$$O = AUH$$




Initial values




Iteration #1




Iteration #2



Iteration #3



Iteration #4



Iteration #5

A note on policies

- A policy over an MDP with primitive actions is a *Markov policy*:


$$\pi: S \times A \rightarrow [0, 1]$$

- A policy over an SMDP with options could also be Markov:


$$\pi: S \times O \rightarrow [0, 1]$$

- But the policy in the original MDP may not be
 - The probability of taking an action at a state depends on the option currently running.


Example



Option A



Option B



Policy

Semi-Markov policies


- A Markov policy for an SMDP may result in a **semi-Markov** policy for the underlying MDP
- (Even if the options are Markov options!)
- *Here, semi-Markov means that the probability of taking a primitive action at each step depends on more than the current state*

Summary

- Original problem: MDP
- MDP + Options = SMDP
- Options framework allows us to both *express a low-level policy*, and plan and learn using the *higher-level SMDP*
- Additionally, the ability to:
 - Create *new* options
 - *Update* option policies
 - *Learn* with options
 - *Interrupt* them ...

What are skills for?

- Adding an option changes the connectivity of the MDP
- This affects:
 - Learning and planning
 - Exploration
 - State-visit distribution
 - Branching factor
 - **Diameter!**



(Sutton, Precup and Singh, AIJ 1999)

Where do options come from?



- Good question!
- Locate **bottleneck** states
 - Learn options to reach these **subgoals** [Simsek and Barto, 2008]
- Could be **extracted** from **solution** to existing tasks
 - NPBRs [Ranchod, Rosman, Konidaris, 2015]
- Options that minimise planning are **NP-hard** [Jinnai, et al, 2019]

Hierarchies of Abstract Machines

(Parr, 1998)


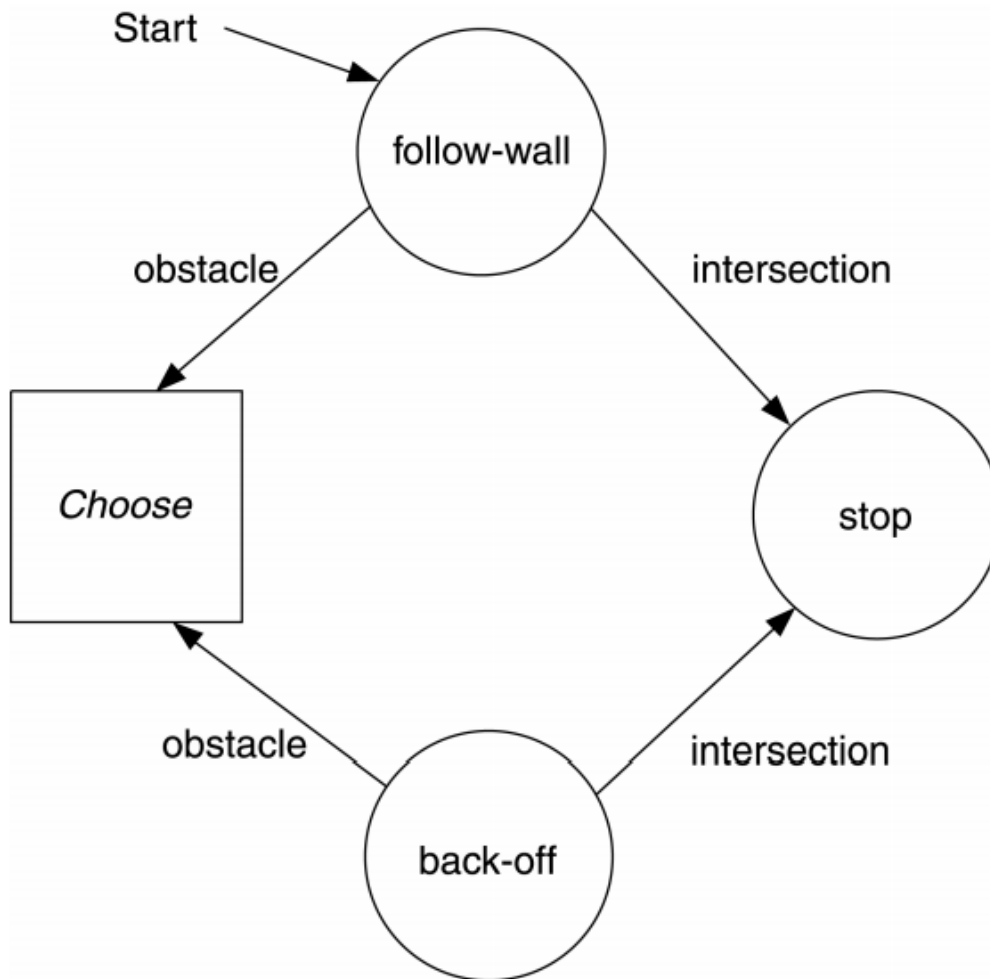
- **Policies** of core MDP are defined as **programs**
- These execute based on **state of MDP**, as well as own **internal state**
- Programs modelled as **finite state machines** (FSMs)
- There are four machine states:
 - **Action** states: execute action in environment
 - **Call** states: execute another FSM as a subroutine
 - **Choice** states: stochastically select a next machine state
 - **Stop** states: halt execution and return control to previous machine

Maze Navigation



Environment

Example of a HAM



HAMs

- Action, call, choice, stop states
- New state space is HAM states \times MDP state space
 - $\mathcal{H} \times S$ state space
- Action, call and stop states are predefined!
 - So **learning** happens only for the **choice** states!
- Can apply **Q-learning to learn choice states** for FSMs


Discussion

- Restrict the set of possible policies
- Good way of injecting prior knowledge
- HAMs + MDP = SMDP
- Link between programming and control
(Programmable HAMS, Stuart and Russell 2001)
- No large scale applications to date


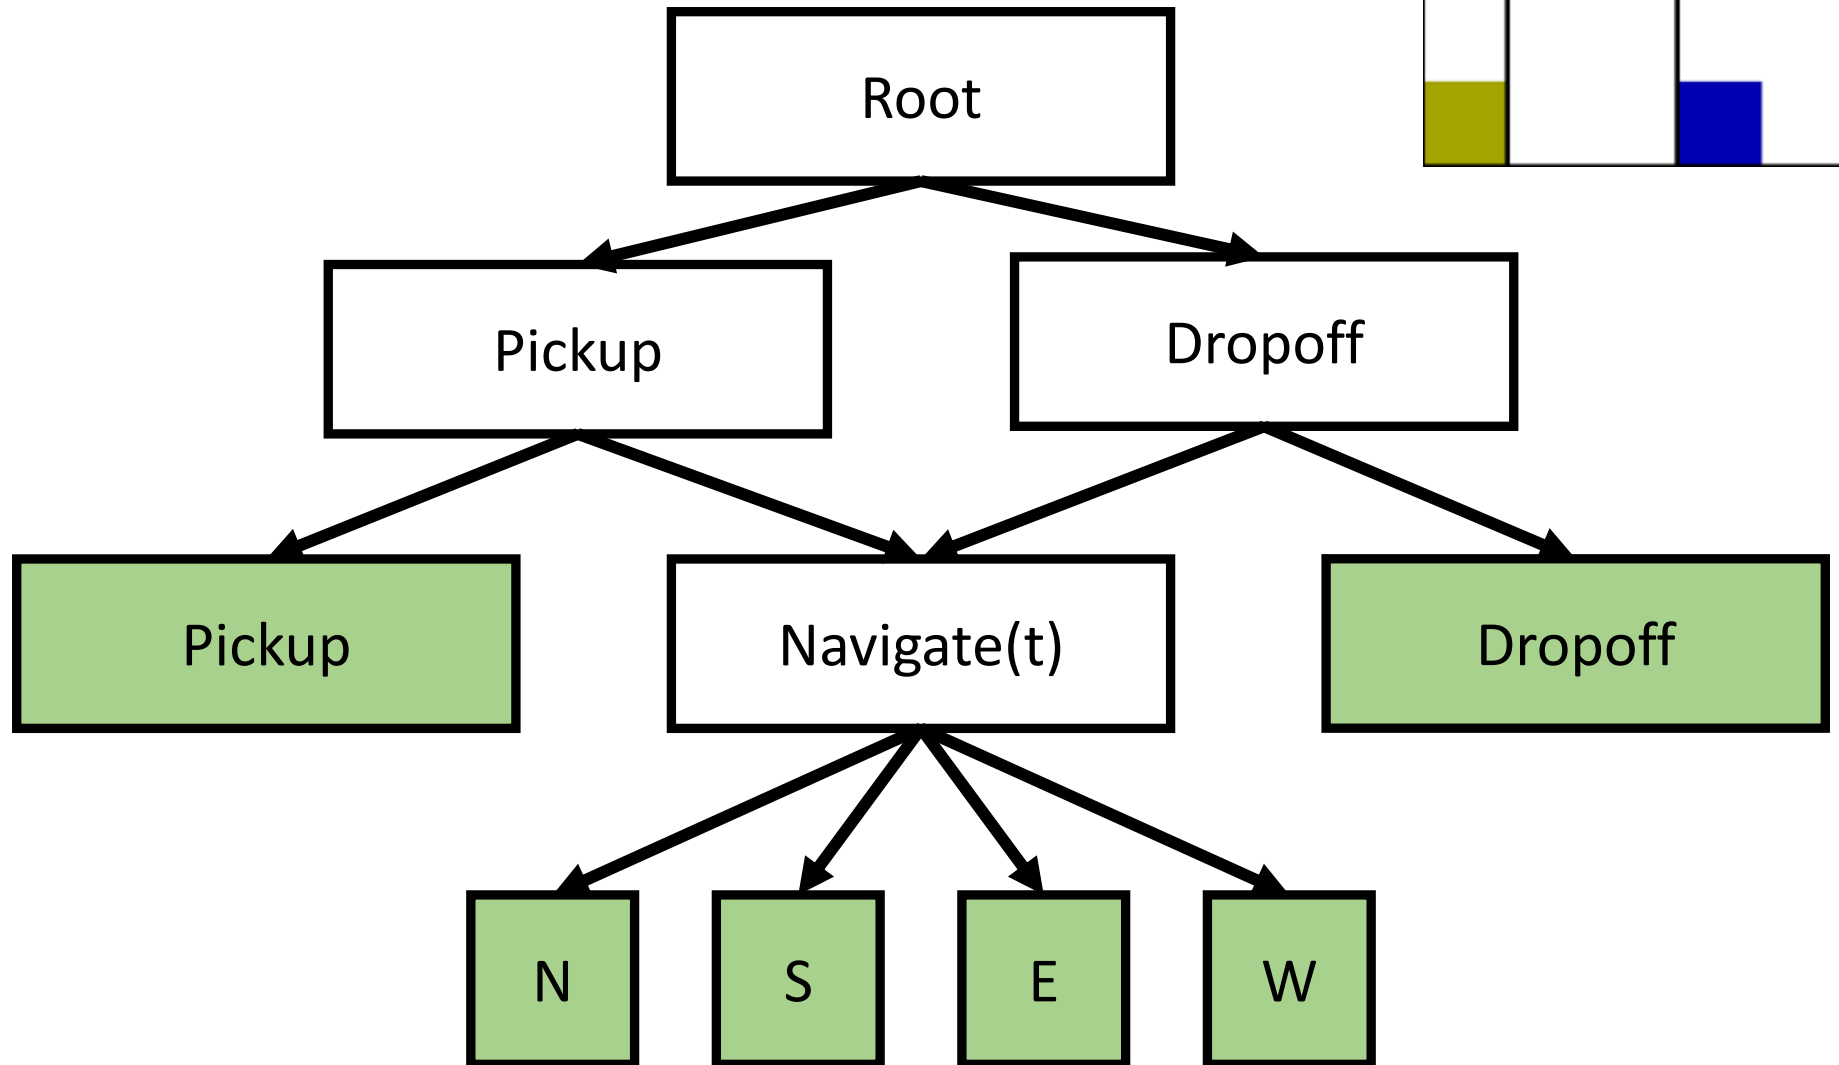
MAXQ value function decomposition

[Dietterich, 2000]


- Taxi domain
 - 5x5 grid. Taxi can move N/S/E/W
 - Collect passenger at one of the colours
 - Drop off at one of the colours
 - -1 for each step, +10 for correct drop-off, -2 for invalid pickup, -5 for wrong dropoff



Decomposition



Formalising



- Let original **base MDP** be $M = (S, A, T, R)$
- Break down into subtasks: M_0, M_1, \dots, M_k
- For each subtask, define $M_i = \{S_i, A_i, \bar{R}_i\}$
 - S_i : set of **states** where subtask is **active**
 - **Terminal states** are all states not in S_i
 - A_i : set of **actions** available
 - **Primitive actions**, and **subtasks** specified by DAG
 - \bar{R}_i : **pseudo-reward** associated with subtask
 - E.g. for Navigate(G), get reward for reaching G

Even if the optimal policy was not to go to G!

Subtasks are SMDPs

- Each M_i is an SMDP with
 - State space S_i
 - Action space A_i
- **Augment** state space with K , a **stack** containing names and params of calling subtasks
 - Like **call stack** in programming languages
- Subtask's policy is **Markov** w.r.t. **augmented state**
- Transition probabilities are well-defined given policies of lower-level subtasks

MAXQ Discussion

- **Real** hierarchical decomposition of a task
- Easy **reuse** of sub-policies
- Very **complex** structure
- Learns **recursively optimal policy**
 - Policy for a parent task is optimal given the learnt policies of its children
 - **Context-free!**
- Recursively optimal policies may be highly **suboptimal** policies

Summary

- RL suffers from **curse of dimensionality**
- HRL seeks to **decompose** problems
- Can **abstract** states/actions
- **Transfer** between tasks
- Focus on **general** problem solving
 - Long lived agents
- Strong AI?!?!?!?!?

Homework

- Complete assignments
- Get 100% on everything
- Apply for MSc and/or PhD
- Publish paper on RL topic 😊
- Build all the AIs!

