


Model-Based RL

Prof. Benjamin Rosman

Benjamin.Rosman1@wits.ac.za / benjros@gmail.com

What is RL?

*An approach that enables an agent to **learn** to act optimally **through trial-and-error interaction** with its environment.*



But this could take forever!

Previously

- Dynamic programming:
 - Given **transition dynamics and reward function**
 - Compute value function/policy
- Model-free RL:
 - Learn a value function/policy from experience
- Model-based RL:
 - **Learn a model** from experience
 - Compute value function/policy

Call this the "model"



Learning and planning

- **Learning**: compute a model/value function/policy based on **real-world experience**
- **Planning**: compute a value function/policy based on **simulated** experience
- Question: What is dynamic programming?


What is dynamic programming?

An interesting question is, 'Where did the name, dynamic programming, come from?' The 1950s were not good years for mathematical research. We had a very interesting gentleman in Washington named Wilson. He was Secretary of Defense, and he actually had a pathological fear and hatred of the word, research. I'm not using the term lightly; I'm using it precisely. His face would suffuse, he would turn red, and he would get violent if people used the term, research, in his presence. You can imagine how he felt, then, about the term, mathematical.

*The RAND Corporation was employed by the Air Force, and the Air Force had Wilson as its boss, essentially. Hence, I felt I had to do something to shield Wilson and the Air Force from the fact that I was really doing mathematics inside the RAND Corporation. What title, what name, could I choose? In the first place I was interested in planning, in decision making, in thinking. But planning, is not a good word for various reasons. I decided therefore to use the word, 'programming.' I wanted to get across the idea that this was **dynamic**, this was multistage, this was time-varying—I thought, let's kill two birds with one stone. Let's take a word that has an absolutely precise meaning, namely dynamic, in the classical physical sense. It also has a very interesting property as an adjective, and that is it's impossible to use the word, dynamic, in a pejorative sense. Try thinking of some combination that will possibly give it a pejorative meaning. It's impossible. Thus, I thought dynamic programming was a good name. It was something not even a Congressman could object to. So I used it as an umbrella for my activities.*

-Richard Bellman

Model-based RL

- General approach:
 - Collect data from the real world
 - Learn a model of the world from the data
 - Use that model to plan
-  *Simulation
"in your head"*
- Problems?
 - Advantages?



Sample efficiency

- Minimise our interaction with the world

```
env = gym.make('MountainCar-v0')
```

```
action = ...
```

```
env.step(action)      Minimise # step!
```



PILCO



Deisenroth, et al. *PILCO: A Model-Based and Data-Efficient Approach to Policy Search*

Models

- **Model**: anything the agent can use to predict how the environment will respond to its actions
- 1. **Distribution model**: description of all possibilities and their probabilities
 - e.g. $p(s', r \mid s, a)$ for all s, a, s', r
- 2. **Sample model**, a.k.a. a simulation model
 - produces sample experiences for given s, a
 - much easier to come by
- Both types of models can be used to produce **hypothetical experience**

Model learning

- Model is an **estimate** of a true MDP $\langle S, A, P, R, \gamma \rangle$
- Assume states and actions are known
- Then model $M_\eta = \langle \hat{P}, \hat{R} \rangle$ where $\hat{P}, \hat{R} \sim P, R$
- We can draw **next state** and **reward samples** from these

Model learning

- Estimate M_η from experience $\{S_1, A_1, R_2, \dots, S_T\}$
- This is supervised learning!

$$\begin{aligned} S_1, A_1 &\rightarrow R_2, S_2 \\ S_2, A_2 &\rightarrow R_3, S_3 \end{aligned}$$

Etc.

- Learning $s, a \rightarrow r$ is a regression problem
- Learning $s, a \rightarrow s'$ is a density estimation problem

Batch Table Lookup Model

- Count visits $N(s, a)$ to each state-action pair

$$\hat{P}(s'|s, a) = \frac{1}{N(s, a)} \sum_{t=1}^T \mathbf{1}(S_t, A_t, S_{t+1} = s, a, s')$$

$$\hat{R}(s, a) = \frac{1}{N(s, a)} \sum_{t=1}^T \mathbf{1}(S_t, A_t = s, a) R_t$$

Online Table Lookup Model

- On receiving experience (s_t, a_t, r_t, s_{t+1}) :

$$\hat{P}(\textcolor{red}{s}_{t+1}|s_t, a_t) = \hat{P}(\textcolor{red}{s}_{t+1}|s_t, a_t) + \alpha(1 - \hat{P}(\textcolor{red}{s}_{t+1}|s_t, a_t))$$

$$\hat{P}(\textcolor{red}{\hat{s}}|s_t, a_t) = \hat{P}(\textcolor{red}{\hat{s}}|s_t, a_t) + \alpha(0 - \hat{P}(\textcolor{red}{\hat{s}}|s_t, a_t))$$

$$\hat{R}(s_t, a_t) = \hat{R}(s_t, a_t) + \alpha(r - \hat{R}(s_t, a_t))$$

Example

- Three states (A, B, C); 2 actions (p, q); 8 episodes
- A, p, 0, B, p, 0, C
- A, q, 0, B, q, 0, C
- B, p, 1, C
- B, p, 1, C
- B, q, 0, C
- B, p, 1, C
- A, p, 0, B, p, 0, C
- A, q, 0, B, p, 1, C

Example

- Three states (A, B, C); 2 actions (p, q); 8 episodes
- A, p, 0, B, p, 0, C
- A, q, 0, B, q, 0, C
- B, p, 1, C
- B, p, 1, C
- B, q, 0, C
- B, p, 1, C
- A, p, 0, B, p, 0, C
- A, q, 0, B, p, 1, C

State-action pair	Transition	Reward
(A, p)	B w.p. 1	0
(A, q)	B w.p. 1	0
(B, p)	C w.p. 1	1 w.p 2/3 0 w.p 1/3
(B, q)	C w.p. 1	0

Sample-based planning

- Generate samples from the model **only!**
 - **Sample** experience from model
- Apply **model-free RL** (MC/TD/etc) to samples

Random-sample one-step tabular Q-planning

Loop forever:

1. Select a state, $S \in \mathcal{S}$, and an action, $A \in \mathcal{A}(S)$, at random
2. Send S, A to a sample model, and obtain
a sample next reward, R , and a sample next state, S'
3. Apply one-step tabular Q-learning to S, A, R, S' :
$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$




*Exact same
Q-learning update!*

Integrating planning and learning

- Model-Free RL
 - No model
 - Learn value function (and/or policy) from real experience
- Model-Based RL (using Sample-Based Planning)
 - Learn a model from real experience
 - Plan value function (and/or policy) from simulated experience
- Dyna
 - Learn a model from real experience
 - Learn and plan value function (and/or policy) from real and simulated experience

Dyna



Dyna

Tabular Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

- (a) $S \leftarrow$ current (nonterminal) state
- (b) $A \leftarrow \varepsilon$ -greedy(S, Q)
- (c) Take action A ; observe resultant reward, R , and state, S'
- (d) $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$
- (e) $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)

(f) Loop repeat n times:

$S \leftarrow$ random previously observed state

$A \leftarrow$ random action previously taken in S

$R, S' \leftarrow Model(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$

*Model-free
Q-learning*


*Model
learning*

Planning!

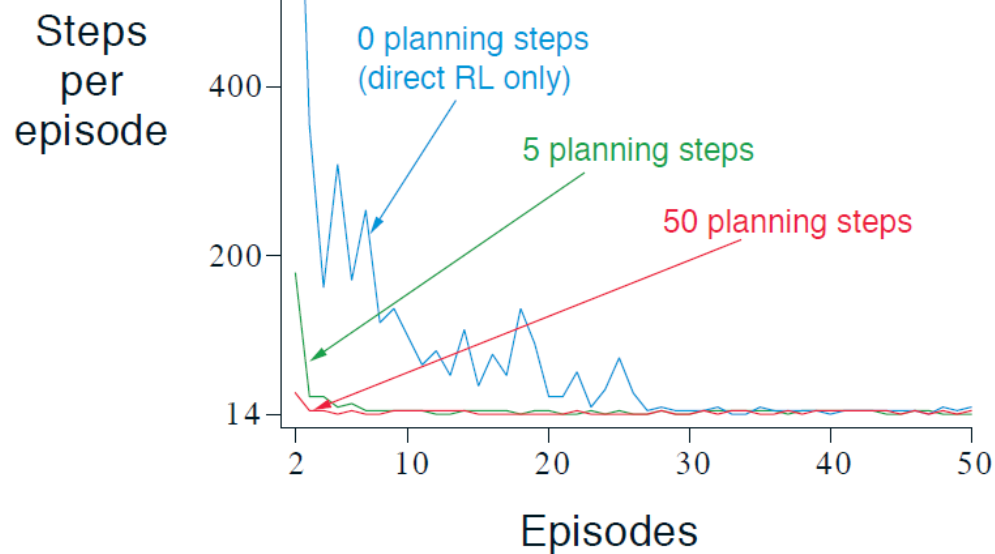



Sample Q-planning

Dyna maze


WITHOUT PLANNING ($n=0$)



WITH PLANNING ($n=50$)



Model error



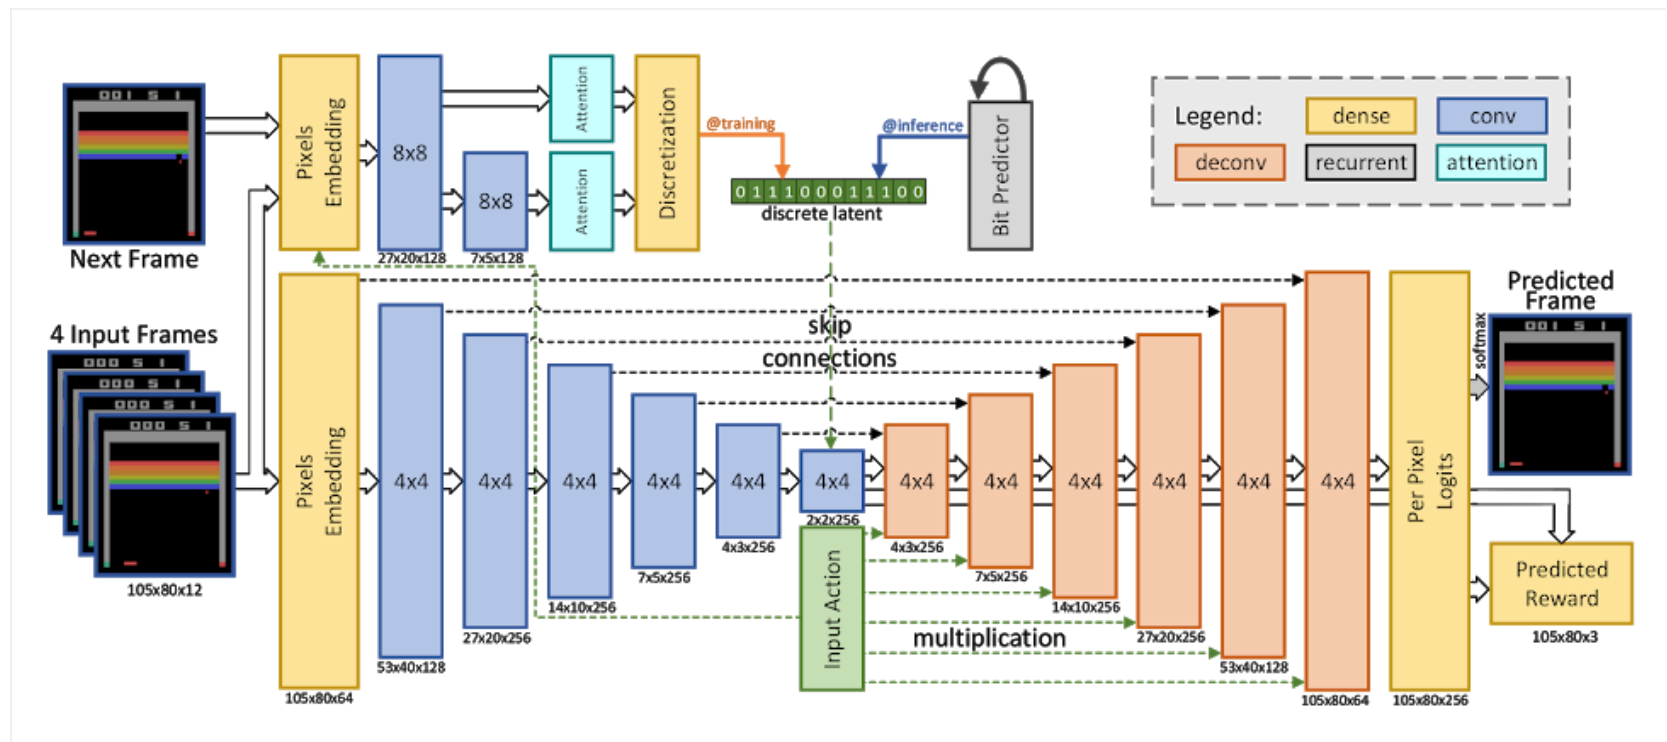
- What if the model is wrong?!
 - Planning would result in **suboptimal** policy
- Why might a model be wrong?
- In tabular case, simple heuristics can be effective

Stochasticity:

Dyna-Q+: give **exploration** reward **bonus** $\kappa\sqrt{\tau(s, a)}$
Time since last visit

- Cutting-edge research for high-dimensional domains (e.g. PlaNet)

SimPle







с: 3, г: 0 с: 7, г: 0



Environment changes



Environment changes



Back to Dyna

Tabular Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

(a) $S \leftarrow$ current (nonterminal) state

(b) $A \leftarrow \varepsilon$ -greedy(S, Q)

(c) Take action A ; observe resultant reward, R , and state, S'

(d) $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$

(e) $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)

(f) Loop repeat n times:

$S \leftarrow$ random previously observed state

$A \leftarrow$ random action previously taken in S

$R, S' \leftarrow Model(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$

*Can we do better than
random?!*

Prioritised sweeping

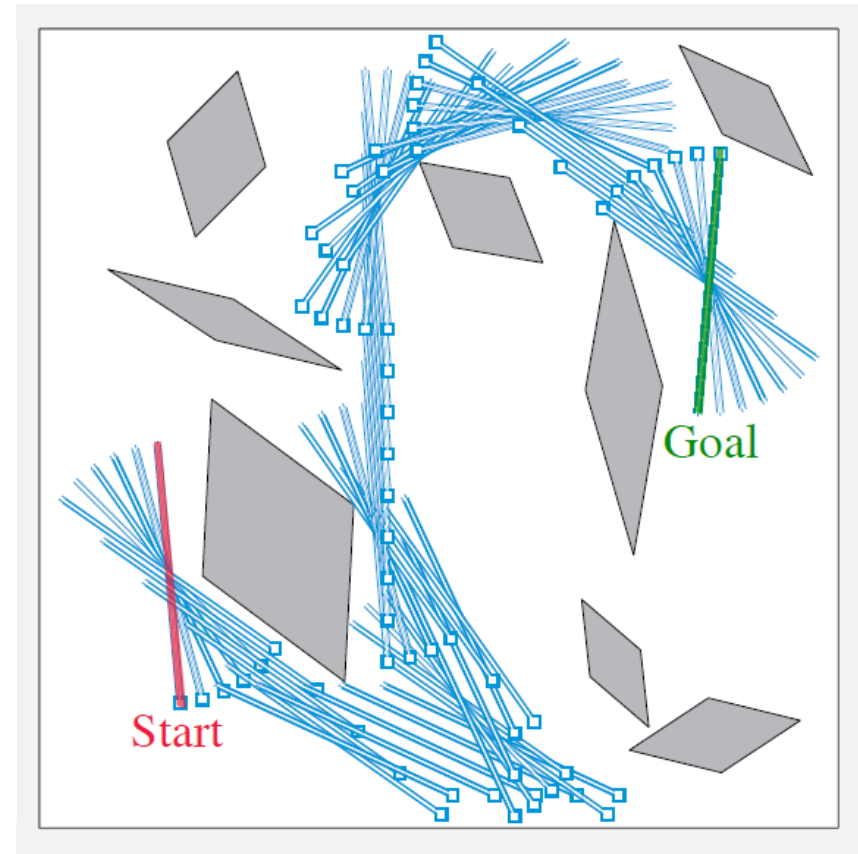

Prioritized sweeping for a deterministic environment

Initialize $Q(s, a)$, $Model(s, a)$, for all s, a , and $PQueue$ to empty

Loop forever:


- (a) $S \leftarrow$ current (nonterminal) state
- (b) $A \leftarrow policy(S, Q)$
- (c) Take action A ; observe resultant reward, R , and state, S'
- (d) $Model(S, A) \leftarrow R, S'$
- (e) $P \leftarrow |R + \gamma \max_a Q(S', a) - Q(S, A)|$. *How big a value change?*
- (f) if $P > \theta$, then insert S, A into $PQueue$ with priority P
- (g) Loop repeat n times, while $PQueue$ is not empty:
 - $S, A \leftarrow first(PQueue)$ *Prioritise bigger changes!*
 - $R, S' \leftarrow Model(S, A)$
 - $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
 - Loop for all \bar{S}, \bar{A} predicted to lead to S :
 - $\bar{R} \leftarrow$ predicted reward for \bar{S}, \bar{A}, S
 - $P \leftarrow |\bar{R} + \gamma \max_a Q(S, a) - Q(\bar{S}, \bar{A})|$. *Update priorities for adjacent states*
 - if $P > \theta$ then insert \bar{S}, \bar{A} into $PQueue$ with priority P

Prioritised sweeping



Monte Carlo Search


- Using a model and a **rollout policy** π_d
 - **Simulate** many episodes with **rollout policy**
 - Compute **mean** return of episodes
 - This is $v_{\pi_d}(s)$ or $q_{\pi_d}(s, a)$
- Then act **greedily**
- By policy improvement theorem, this is **better** or equal to π_d
- Turns out to be quite **effective**!



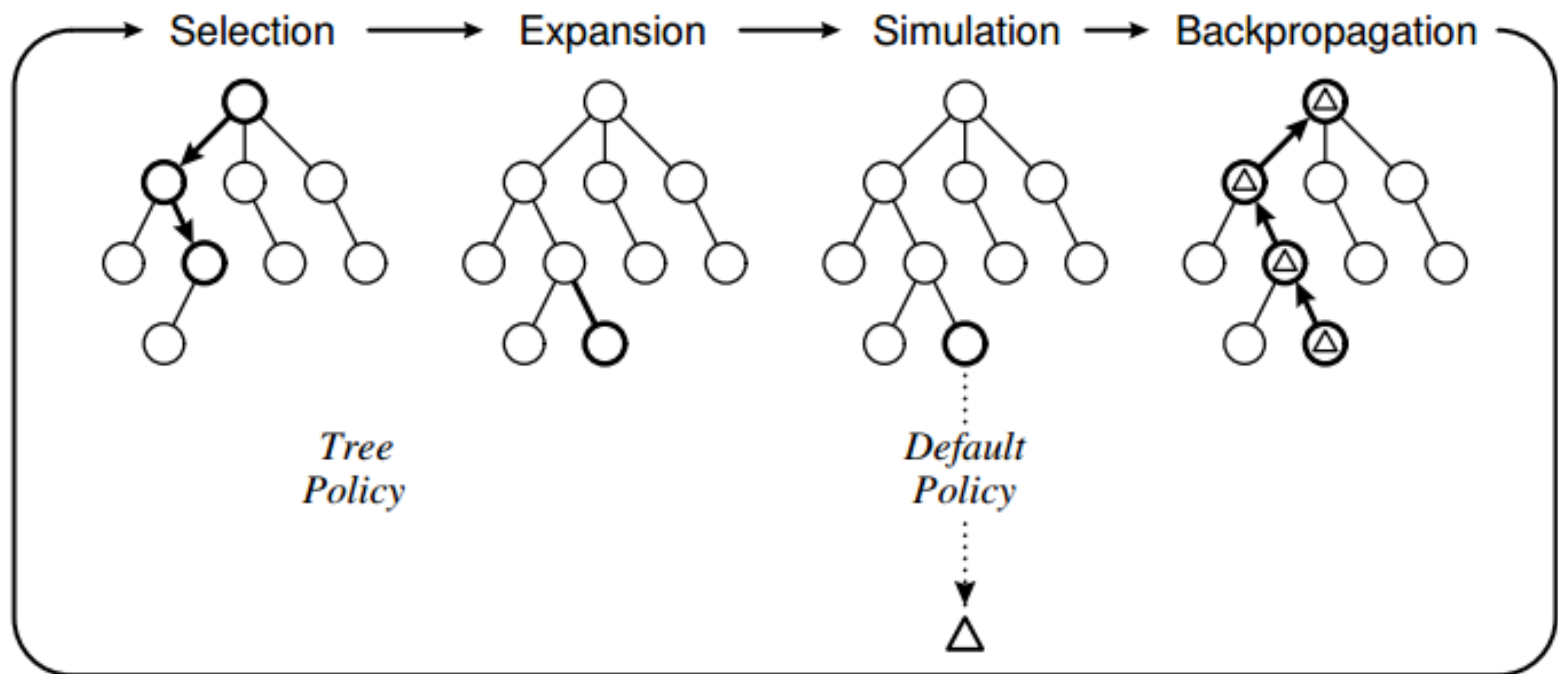
Monte Carlo Tree Search (MCTS)

- A key planning algorithm in RL
- Combine Monte Carlo search with bandit theory (remember UCB?)
- At each state (node), store score and visit count
- Build an asymmetric search tree
- Visits most promising states more often
- Can stop planning and return action anytime

MCTS



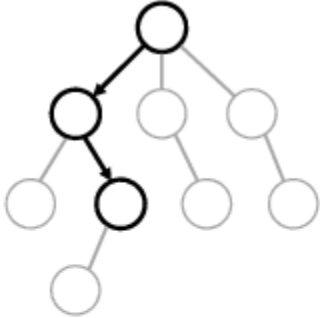
MCTS Phases



Selection

- Start at **root node** (current state)
- Use **tree policy** to descend tree to leaf node
- **UCB1** selection: at each node s , select child i maximising

Value of node $\rightarrow X_i + C_p \sqrt{\frac{\ln(n_s)}{n_i}}$ *Visits of parent* $\rightarrow \ln(n_s)$ *Visits of node* $\rightarrow n_i$




- MCTS + UCB1 = UCT

(a) Selection

Expansion


- At leaf node, take **unexplored** action
- **Add new** node to tree



(b) Expansion

Simulation


- From new node, use rollout policy to simulate trajectory
- At end of episode, get reward



(c) Simulation

Backpropagation

- **Update** all nodes from leaf to root with **reward**
- **Update** all nodes' **visit** counts by 1




(d) Backpropagation

MCTS

- Continue cycle of selection, expansion, simulation and backprop until **time runs out**
- Select **best action**
 - Action that leads to state with **highest value** (max child)
 - Action that leads to state with **most visits** (robust child)
- Can **reuse tree** for next action!

MCTS for Go



Summary

- Distinction between **learning** and **planning**
- Use experience from real world to **learn model**
- Use model to update policy/value function
 - Improve **sample efficiency**
- Model error is an issue
 - Hard to learn models for complex environments
 - But we do it!

Homework

1. Think about your assignment
2. Implement your thoughts in the aforementioned step
3. Train your agent building on steps 1 and 2
4. Submit your assignment