

Reinforcement Learning – COMS4061A/7071A

Model-Free Learning



Prof. Benjamin Rosman

Benjamin.Rosman1@wits.ac.za / benjros@gmail.com

Last time


- MDP: $\langle S, A, P, R, \gamma \rangle$
 - Known dynamics and reward function: $\langle S, A, \textcolor{red}{P}, \textcolor{red}{R}, \gamma \rangle$

$$v_{k+1}(s) = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_k(s')]$$



This week

- The full RL problem
 - We know nothing!



Challenges?

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s', r | s, \pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)

We can't just query for every possible state!

3. Policy Improvement

policy-stable \leftarrow true

For each $s \in \mathcal{S}$:

old-action $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$

If *old-action* $\neq \pi(s)$, then *policy-stable* \leftarrow false

If *policy-stable*, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

One-step look-ahead needed!

MONTE CARLO METHODS

Prediction and Control

Dr Strange uses Monte Carlo



14 000 605 “episodes”; 1 win

If +1 for winning and 0 otherwise

$$v_{\pi}(s) = \frac{1}{14000605}$$

Blackjack Example

- Previously, to compute value function, we needed the dynamics
- But computing $p(s', r | s, a)$ is **non-trivial**
 - e.g. I have 13, what's the probability I get to 15 given the cards on the table?
- But MC is easy!
 - Just **play out** many games (according to policy) from that position and **average**!

Monte Carlo prediction

- Compute $v_{\pi}(s)$
- For **episodic tasks**, the return is:

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots \gamma^{T+1} R_T$$

- The value function is the **expected return**

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t \mid S_t = s]$$

- Start at s , execute episodes according to π and take average
 - **Empirical mean**

Monte Carlo prediction

First-visit MC prediction, for estimating $V \approx v_\pi$

Input: a policy π to be evaluated

Initialize:

$V(s) \in \mathbb{R}$, arbitrarily, for all $s \in \mathcal{S}$

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Loop forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless S_t appears in S_0, S_1, \dots, S_{t-1} :

Append G to $Returns(S_t)$

$V(S_t) \leftarrow \text{average}(Returns(S_t))$

*For "every-visit",
delete this check*

Note: no dependence on $|\mathcal{S}|$

Incremental mean

- For each state S_t with return G_t

$$\begin{aligned} N(S_t) &\leftarrow N(S_t) + 1 \\ V(S_t) &\leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t)) \end{aligned}$$

- In **non-stationary problems**, we may want to **forget old episodes**:

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

Monte Carlo control

- We can evaluate π
- So we just need **policy improvement**
 - Then we can do **policy iteration**!
- But how do we improve the policy with $v_\pi(s)$?
- We instead estimate $q_\pi(s, a)$ and take **argmax**
- Must ensure all state-action pairs are visited an infinite number of times (**exploring starts**) for **convergence**

Monte Carlo Exploring Starts

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$


Initialize:

$\pi(s) \in \mathcal{A}(s)$ (arbitrarily), for all $s \in \mathcal{S}$

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

*Explore all states
infinitely by random
starts*



Loop forever (for each episode):

Choose $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$ randomly such that all pairs have probability > 0

Generate an episode from S_0, A_0 , following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$\pi(S_t) \leftarrow \arg\max_a Q(S_t, a)$

Exploration!

Different from argmax!



- ε -soft policies: **any** policy where $\pi(a | s) > 0$ for all states, actions
- With probability ε , pick suboptimal action at random; otherwise, greedy action

$$\pi(a | s) = \begin{cases} 1 - \varepsilon + \frac{\varepsilon}{|A|}, & \text{if } a = \operatorname{argmax} Q(s, \cdot) \\ \frac{\varepsilon}{|A|}, & \text{otherwise} \end{cases}$$

Shortcomings of Monte Carlo

- Need episodes!
 - i.e. works in the **episodic** case only
 - Must have policies that can actually **terminate** the episode
 - Must wait until the end of the episode until we can update
- High **variance**
 - The return depends on many random actions, transitions and rewards

If one had to identify one idea as central and novel to reinforcement learning, it would undoubtedly be temporal-difference (TD) learning. TD learning is a combination of Monte Carlo ideas and dynamic programming (DP) ideas.

– Sutton and Barto [2018]

TEMPORAL DIFFERENCE

Prediction

Temporal difference learning

- TD learns from experience
- Model-free
- Learns from **incomplete** episodes
 - Can learn **without** the final outcome
 - And also in the **continuing** case
- Learns by **bootstrapping**
 - Similar to DP

TD Update

- MC update:

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

*Update toward full
episode return*



- TD(0) update:

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

*Update toward estimated
return*



- **TD target:** $R_{t+1} + \gamma V(S_{t+1})$
- **TD error:** $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$

Policy evaluation:

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'


$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal


Dynamic Programming

$$V(S_t) \leftarrow \mathbb{E}_{\pi}[R_{t+1} + \gamma V(S_{t+1})] = \sum_a \pi(a | s) \sum_{s', r} p(s', r | S_t, a) [r + \gamma V(s')]$$




Monte Carlo

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$$



TD(0)

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$



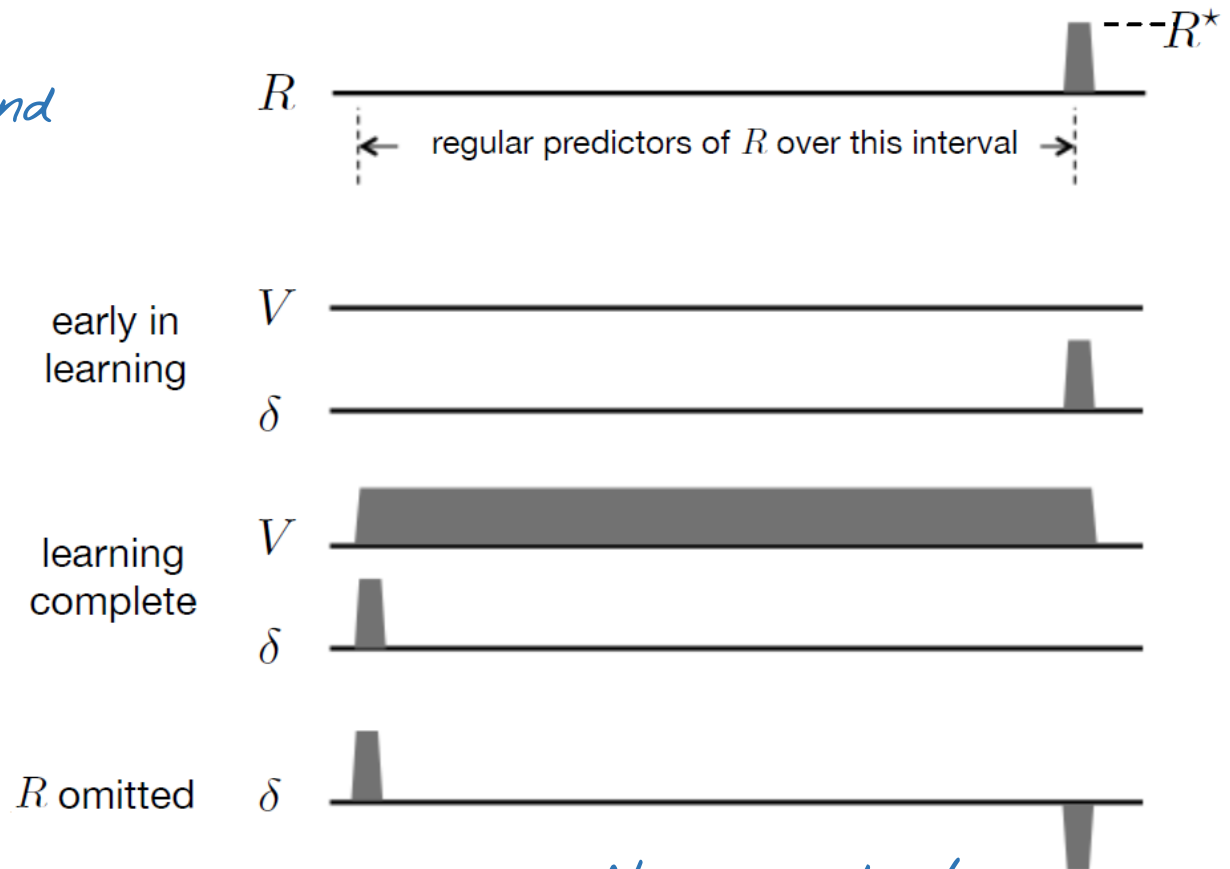
Aside

- $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is the **learning signal**
- $V(S_t) \leftarrow V(S_t) + \alpha \delta_t$
- The signal is high/low when something **unexpectedly** positive/negative happens

No discounting
Single reward at end

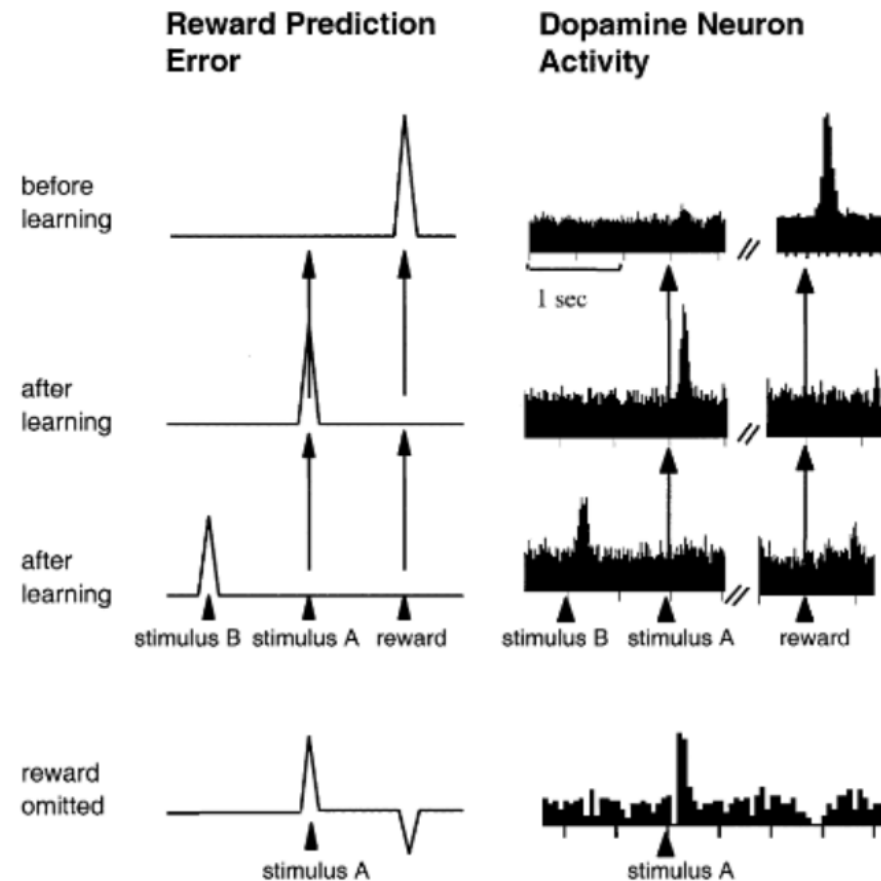
V is zero initially.
Then we see a reward!

For any state
that enters the
predictive states



No reward when
it was expected!
Negative error

Dopamine



TEMPORAL DIFFERENCE

Control

TD Control

- We have **no model**
 - Learning **V is problematic**. Why?
- We will use TD learning to update the **Q -value function**.
- Previously, we used $(S_t, A_t, R_{t+1}, S_{t+1})$
- Now we use $(S_t, A_t, R_{t+1}, S_{t+1}, \mathbf{A}_{t+1})$

TD Control

- 2 methods:
 - SARSA
 - Q-learning
- Difference is how they select A_{t+1}
- SARSA: select A_{t+1} according to policy that selected A_t (e.g. with ϵ -greedy exploration)

$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

- Q-learning: $A_{t+1} = \operatorname{argmax}_a Q(S_{t+1}, a)$

$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Note on convergence

- Greedy in the limit with infinite exploration (**GLIE**)
 - All state-action pairs visited **infinitely many times**
 - Policy **converges to greedy policy** e.g. ε -greedy with $\varepsilon_k = \frac{1}{k}$
- Robbins-Monro **sequence of learning rates** α_t :

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

SARSA

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Same policy

Q-learning

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

 *not ε -greedy*

On-policy vs off-policy

- On-policy methods learn by following the current policy
- Off-policy methods learn from a different **behaviour policy**
- SARSA is on-policy. Q-learning is off-policy
 - Why?
- Why is this important?
 - Learn from **observing humans** or other **agents**
 - **Re-use experience** generated from **old policies** (e.g. replay buffers)
 - Learn about optimal policy while following **exploratory policy**

N-STEP RETURNS



SARSA(λ)

Between TD and MC

MC Methods (∞ -step lookahead)

TD(0) Methods (one-step lookahead)

WHAT ABOUT 2-steps ahead? 3? 4? ...



n-step returns


- $n = 1$ (TD) $G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$
- $n = 2$ $G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2})$
- ...
- $n = \infty$ (MC) $G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^n R_T$
- n-step TD learning:

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^{(n)} - V(S_t))$$


n-step returns

- More computation required
- But better estimates:
 - n -step return better than $(n - 1)$ -step return for approximating v_π


Path taken



Action values increased by one-step Sarsa



Action values increased
by 10-step Sarsa




λ -returns

- How about **averaging** n -step returns over different n ?

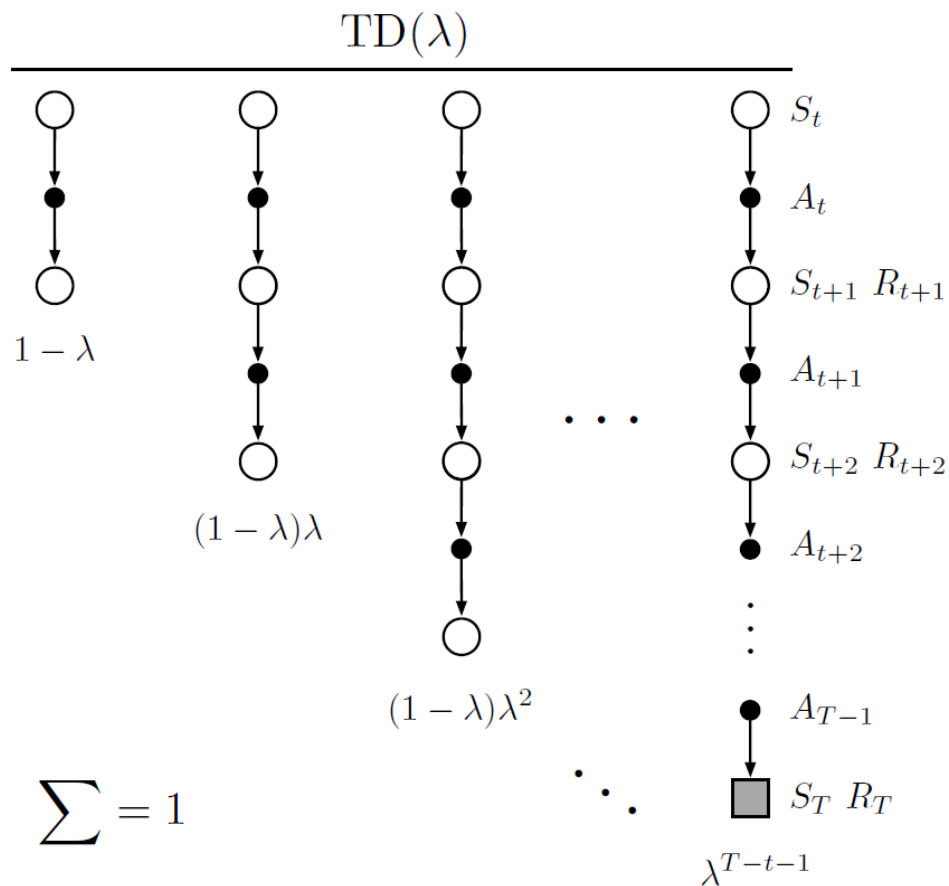
$$\frac{1}{2} G_t^{(1)} + \frac{1}{2} G_t^{(2)}$$

- How about averaging over **all** n ?
- The λ -return combines all n -step returns with weights $(1 - \lambda)\lambda^{n-1}$


$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

TD target

λ -returns




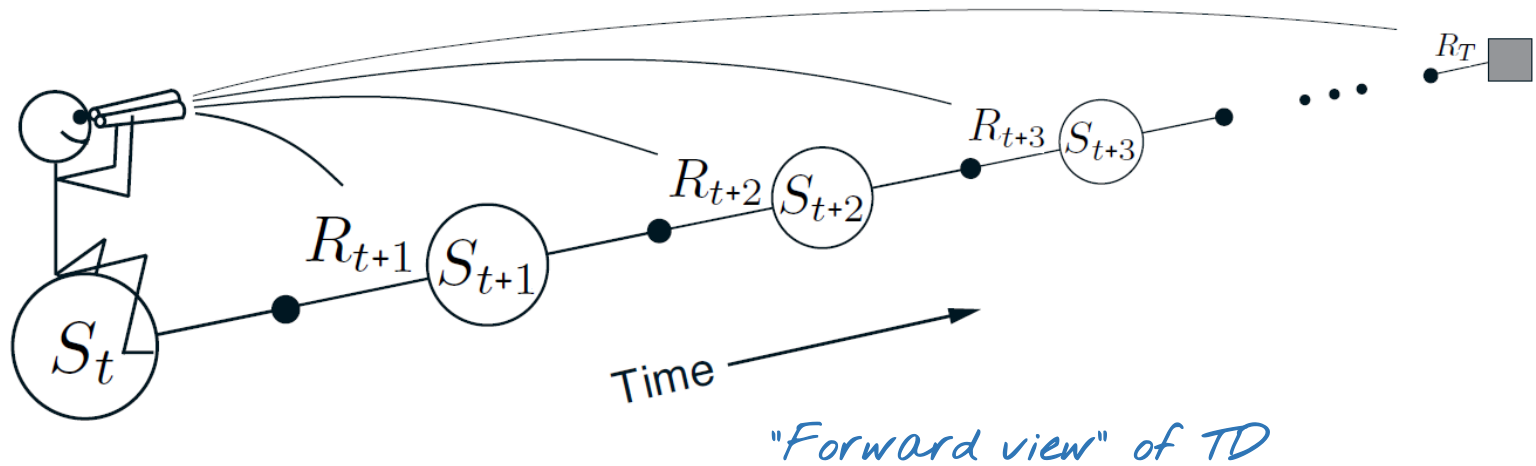


Figure 12.2: Weighting given in the λ -return to each of the n -step returns.

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

- What happens if $\lambda = 0$?
 - And for $\lambda = 1$?

n-step and λ -returns




- Need to have **knowledge** of **future** rewards
- But we only know the **present**
- So we must wait until the **end of the episode** 😞

Eligibility traces

- A way of **assigning credit backward** in time
- **Frequency heuristic**: assign credit to most frequent states
- **Recency heuristic**: assign credit to most recent states
- Eligibility traces combine both heuristics

$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$




TD(λ)

- Keep an eligibility trace for every state s
- Update value in proportion to **TD error** and **eligibility trace**

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

*Roughly equivalent to
forward view*



TD(λ) for policy evaluation

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

Initialize S

$e \leftarrow 0$

(a vector the size of S)

Loop for each step of episode:

$A \leftarrow$ action given by π for S

Take action A , observe R, S'

$\delta = R + \gamma V(S') - V(S)$

$e(S) \leftarrow e(S) + 1$

For every state $t \in \mathcal{S}$:

$V(t) \leftarrow V(t) + \alpha \delta e(t)$

$e(t) \leftarrow \gamma \lambda e(t)$

*Note when $\lambda = 0$ we get
TD(0)*

(decay the eligibility trace of t)

$S \leftarrow S'$

until S is terminal

Summary

- Methods that learn from experience without a model

- MC: Update toward the **full return**:


$$\delta_t^{MC} = G_t - V(S_t)$$

- TD(0): Update toward **one-step** difference:

$$\delta_t^{TD} = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

- SARSA (**on-policy** TD) vs Q-learning (**off policy** TD)
- TD(λ): **interpolate** between TD and MC methods

Unified View of RL



Homework

- Use the CliffWalking domain from OpenAI gym
 - See Example 6.6, pg 132 in Sutton and Barto [2018]
- Modify the $TD(\lambda)$ algorithm presented to implement $SARSA(\lambda)$
 - The only difference here is that there is an eligibility trace for each **state-action** pair!
 - Use ε -greedy policies with $\varepsilon = 0.1$ and a learning rate of $\alpha = 0.5$
 - Run $SARSA(\lambda)$ on the domain for $\lambda = \{0, 0.3, 0.5\}$ for 200 episodes
 - Record the return for each episode
 - Average your returns over 100 runs

By next week's lecture, submit on Moodle:

1. Perform a single run of the algorithm. After each episode plot the value function (take $\max_a Q(s, a)$) learned so far as a heatmap for each λ side by side. This should^a result in 200 separate plots/images. Turn these images into an animation/video and submit it.
2. A combined plot of average return over time for the different values of λ . Include error bars/shading indicating variance in your results
3. Your code